BGS 270: Law, Ethics, and A.I. Syllabus

Professor David Atkinson
Spring 2023
Room: SZB 4.414

NOTE: The content of this syllabus is subject to change based on new developments in AI and the pace of classroom discussions. Do not read more than one week in advance without double-checking with me

Learning Content

1. Each class will cover one area of A.I.
2. There will be two kinds of assignments for each class: the *main* assignments and then the *supplementary* assignments.
    1. Examples of main assignments:
        1. News Articles
        2. Law review articles
            1. Don't be fooled by the number of total pages. Law review articles have extra wide margins and you don't need to read the extensive footnotes, so often there are only 300 or fewer words per page.
    2. Examples of supplementary assignments:
        1. TED Talks
        2. YouTube videos
        3. Documentaries
        4. Podcast episodes
        5. Non-scholarly articles
    3. For each class you will do *all* the main assignments and a specified number of supplementary assignments (usually just one).

Class Structure

1. Magnets:
    1. I'll put everyone's name on a magnet, and I'll put all the magnets on the board in a cluster on the left
    2. As you enter class, find your magnet, then move it to the supplemental assignment you chose for that class, which will be on the right side of the board
    3. This lets me know who read what and it serves as a way for me to take attendance

2. We'll then go into the new material from the main assignments

1. For the first third of class:
    1. I'll introduce the day's topic and lead discussion of the main assignment, which everyone should have read.
    2. We'll discuss the assignment/topic as a class Socratically/conversationally
2. For the remaining thirds of class:
    1. I will randomly choose one of the supplementary groups to discuss
    2. If nobody speaks up, I will use a random number generator to select a person to introduce the topic from their supplemental assignment group

Grading

1. Persuasive Memos = 35% of the grade
    1. You will write three persuasive memos over the semester.
    2. The due dates for the memos are detailed under the "Classes" section below.
    3. You must turn in a memo no later than the Due Date. You may turn it in earlier if you'd like.
    4. The topic of the persuasive memo must cover one of the topics discussed between when the previous memo was due and when the current memo is due. For example, memo two must be on Social Media and Entertainment, Credit Scores and Government Quantifying, or Job Hiring and the Workforce.
    5. The memo must include the following:
        1. Introduction and summary of the topic as it has existed
        2. A discussion of how AI is impacting the topic
        3. An explainer of the ethical questions and implications
        4. A persuasive argument of what approach society should take to the issue based on the Veil of Ignorance, including an explanation of why your position is the best approach.
        5. Only materials from this syllabus may be cited. The memos must include citations to support your statements.
    6. The memos must be between 1,700 and 2,000 words. Three points will be deducted from your grade for going above or below the word count.
    7. The overall tone of the memos must be persuasive, not objective.
    8. Citation can be simple footnotes that merely refer to the specific syllabus content. Include page numbers if applicable. It does not need to be in MLA or APA format.
    9. The highest score possible per memo is 50 points.
    10. If you have an innovative solution proposal to address problems you identify, feel free to share it. Bonus points may be awarded for clever proposals at my sole discretion.
    11. Grading will be based on the following rubric:

	10 points	8 Points	6 Points	4 Points	2 Points

Critical Analysis	• Excels in responding to assignment, and demonstrates mastery of course concepts and materials • Thesis presents a clear, focused, and compelling argument • Paper recognizes the complexities of its argument throughout the analysis	• Responds appropriately to the assignment, demonstrates clear understanding of course concepts and materials • Good argument, clearly articulated in thesis, though might need refining • Begins to acknowledge the complexities of its argument	• Doesn't fully respond to the assignment, demonstrates some misunderstanding of course concepts and materials • Paper has a weak argument, thesis is too general • Fails to acknowledge other views	• Doesn't respond appropriately to the assignment, disconnected from course concepts and materials • Argument is unclear, thesis is weak • Thesis too vague or general to be nuanced or complicated	• Does not respond to the assignment, displays no familiarity with course concepts or materials • No identifiable argument or thesis
Evidence and Support	• Argument is thoroughly supported by strong, specific, and appropriate evidence • Evidence is clearly introduced, analyzed and connected to the argument	• Paper's argument is supported by relevant evidence, though not always the strongest or specific quotations • Analysis of evidence needs further development	• Paper's argument is supported by limited evidence that is only occasionally relevant • Connections between argument and evidence are somewhat unclear	• Evidence is insufficient, misconstrued or misrepresented • Unclear connections between evidence and argument	• Argument is based on little to no evidence • Connections between evidence and argument are absent/incorrect
Structure	• Paper flows logically to craft a cohesive argument • Paragraphs clearly guide the reader through a progression of ideas • Uses transitional sentences to develop strong relationships between ideas	• Generally well-constructed flow of ideas Paragraphs are ordered thoughtfully, each paragraph relates to central argument • Transitional sentences create a logical progression of ideas	• Paper jumps from one idea to the next, lacking a clear structure • Occasional connection of ideas between paragraphs • Simple sequential rather than transitions based on logic	• Paper wanders from one idea to the next, making it difficult to distill the argument • Limited connection of ideas between paragraphs • Paragraphs may lack topic sentences or connection of ideas	• Lacking organization and coherence • No connection of ideas between paragraphs • Disjointed connection of ideas between paragraphs
Style	• Displays a unique critical voice • Style fits the paper's audience • Chooses words	• Displays a clear critical voice • Style is conscious of paper's audience	• Displays a critical voice that is generic or bland • Style only occasionally	• Critical voice is unclear • Style isn't appropriate for paper's audience	• Lacking critical voice • Unaware of paper's audience

	carefully, for their precise meaning	• Uses words effectively, if too generally at times	displays awareness of paper's audience • Sentence structure and word choice frequently too unfocused, wordy or confusing	• Simple, awkward, or monotonous sentence structure and word choices	• Many awkward sentences and misused words
Mechanics and Citations	• Almost entirely free of spelling, grammar, and punctuation errors • Cites all sources	• May contain a few spelling, grammar, or punctuation errors, but they don't impede understanding • Cites all sources	• Several spelling, grammar, or punctuation errors that distract the reader • Cites most sources	• Contains many spelling, grammar, or punctuation errors • Cites most sources	• Pervasive spelling, grammar, or punctuation errors • Cites half or fewer sources

2. Persuasive Op-ed/Letter = 15% of the grade
    1. You will pick one company, think tank, or newspaper to write to.
    2. The topic must be related to the intersection of A.I. and ethics and must be about a topic we discuss in class.
    3. Your goal will be to persuade the recipient or readers to undertake specific actions for clear ethical reasons
    4. The op-ed/letter must be between 700 and 1,000 words.
    5. You must forward to me proof that you've sent your writing to the intended recipient.

3. Participation = 50% of the grade
    1. Everyone starts with 100% on the participation grade.
        1. Each class is worth 1 participation point, meaning there are only 14 participation points in total.
        2. You lose a full point for missing class.
        3. You lose a half point for choosing to not speak when called on by the Random Number Generator or for being unprepared or unable to demonstrate you read the material.
    2. You can regain a half point by making a great/interesting point in class, asking a great question, being articulate and persuasive, etc. This is awarded at my sole discretion. It will be rare.

4. Late Assignments and Missed Classes
    1. If an assignment is late, but submitted before the subsequent class starting time, you can earn up to half of the total possible points on that assignment.

    2. If the late assignment is submitted after the subsequent class's starting time, you get a 0 for that assignment.

5. Anonymity
    1. All grading for work submitted by the deadline will be graded anonymously.

Learning Outcomes

1. Understand the legal and ethical implications of artificial intelligence and how they impact society, organizations, and individuals.
2. Analyze the potential consequences of AI applications and technologies on various stakeholders, including government, industry, and the general public.
3. Develop a critical understanding of the potential risks and benefits of AI and how to mitigate the risks while maximizing the benefits.
4. Understand the role of law, regulation, and ethics in shaping the development and use of AI.
5. Understand how to apply ethical principles and frameworks to AI-related decision making.
6. Understand the social, cultural, and political context in which AI operates and how these contexts shape the development and deployment of AI.
7. Understand the role of public policy in regulating and shaping the development and use of AI.
8. Understand the role of research ethics in AI and the ethical responsibilities of AI researchers.
9. Understand the role of diverse perspectives and voices in shaping the development and use of AI, including the perspectives of underrepresented groups and marginalized communities.
10. Understand the role of AI in advancing or hindering social justice and equality.

Studying Recommendations

1. Form a study group of 2-4 people.
2. Read the assignments independently and slowly.
3. Take handwritten notes.
4. Then, meet with your group and discuss them. Debate them.

Electronics

1. Generally, no electronics are allowed during class. This means no laptops, no phones, no tablets, no headphones. However, please bring your laptops for some activities.
2. If you need special accommodations, please let me know and we can discuss.

Accessible/Compliant Statement

If you are a student with a disability, or think you may have a disability, and need accommodations please contact Disability and Access (D&A). You may refer to D&A's website for contact and more information: http://diversity.utexas.edu/disability/. If you are already registered with D&A, please deliver your Accommodation Letter to me as early as possible in the semester so we can discuss your approved accommodations.

Accessible, Inclusive, and Compliant Statement

The university is committed to creating an accessible and inclusive learning environment consistent with university policy and federal and state law. Please let me know if you experience any barriers to learning so I can work with you to ensure you have equal opportunity to participate fully in this course. If you are a student with a disability, or think you may have a disability, and need accommodations please contact Disability and Access (D&A). Please refer to D&A's website for contact and more information: http://diversity.utexas.edu/disability/. If you are already registered with D&A , please deliver your Accommodation Letter to me as early as possible in the semester so we can discuss your approved accommodations and needs in this course.

Academic Integrity

Students who violate University rules on academic misconduct are subject to the student conduct process and potential disciplinary action. A student found responsible for academic misconduct may be assigned both a status sanction and a grade impact for the course. The grade impact could range from a zero on the assignment in question up to a failing grade in the course. A status sanction can range from probation, deferred suspension and/or dismissal from the University. To learn more about academic integrity standards, tips for avoiding a potential academic misconduct violation, and the overall conduct process, please visit the Student Conduct and Academic Integrity website at: http://deanofstudents.utexas.edu/conduct.

Classes

Week	Date	Class Topic	Assignments Due
1	1/9/23	What is A.I. and What are Ethics?	
2	1/23/23	Image Recognition	
3	1/30/23	Text Recognition	

4	2/6/23	Speech Recognition	Essay 1 Due
5	2/13/23	Social Media and Entertainment	
6	2/20/23	Credit Scores and Government Quantifying	Survey for feedback
7	2/27/23	Job Hiring and the Workforce	Essay 2 Due
8	3/6/23	Finance and Wall Street	
9	3/20/23	Military and War	
10	3/27/23	• Autonomous Vehicles (guest speaker) • Discuss the op-ed assignment	Essay 3 Due
11	4/3/23	Medicine and Medical Care	
12	4/10/23	The Judiciary	Submit who you will send your op-ed to
13	4/17/23	Board Directors / Corporate Governance	
14	4/24/23	Artificial General Intelligence	Submit proof op-ed or letter was sent

Class 1: What is A.I. and What are Ethics?

- **Main Assignments: Do Them All**
  - The Fairness Principle: The Fairness Principle: How the Veil of Ignorance Helps Test Fairness - Farnam Street
  - Big Data requires human labeling:
    - AI Is Learning From Humans. Many Humans. - The New York Times
  - If AI Is Predicting Your Future, Are You Still Free? | WIRED
  - TED Ed: How does artificial intelligence learn? - Briana Brownell (4:57)
  - Veil of Ignorance - Ethics Unwrapped
  - Ethical dilemma: The burger murders - George Siedel and Christine Ladwig (5:46)
- **Supplementary Assignment: Complete One Option**
  - **Option 1:**
    - Zeynep Tufekci: Machine intelligence makes human morals more important | TED Talk (17:33 Machine intelligence is here, and we're already using it to make subjective decisions. But the complex way AI grows and improves makes it hard to understand and even harder to

control. In this cautionary talk, techno-sociologist Zeynep Tufekci explains how intelligent machines can fail in ways that don't fit human error patterns -- and in ways we won't expect or be prepared for.)

- Mainak Mazumdar: How bad data keeps us from good AI | TED Talk (10:00 The future economy won't be built by people and factories, but by algorithms and artificial intelligence, says data scientist Mainak Mazumdar. But what happens when these algorithms get trained on biased data? Drawing on examples from Shanghai to New York City, Mazumdar shows how less-than-quality data leads to AI that makes wrong decisions and predictions -- and reveals three infrastructural resets needed to make ethical AI possible.)
- Margaret Mitchell: How we can build AI to help humans, not hurt us | TED Talk (9:47 As a research scientist at Google, Margaret Mitchell helps develop computers that can communicate about what they see and understand. She tells a cautionary tale about the gaps, blind spots and biases we subconsciously encode into AI -- and asks us to consider what the technology we create today will mean for tomorrow.)
- You can only save one— who do you choose? - Doug MacKay (4:26)

- **Option 2:**
  - Genevieve Bell: 6 big ethical questions about the future of AI | TED Talk (14:39 Artificial intelligence is all around us ... and the future will only bring more of it. How can we ensure the AI systems we build are responsible, safe and sustainable? Ethical AI expert Genevieve Bell shares six framing questions to broaden our understanding of future technology -- and create the next generation of critical thinkers and doers.)
  - Kriti Sharma: How to keep human bias out of AI | TED Talk (12:01 AI algorithms make important decisions about you all the time -- like how much you should pay for car insurance or whether or not you get that job interview. But what happens when these machines are built with human bias coded into their systems? Technologist Kriti Sharma explores how the lack of diversity in tech is creeping into our AI, offering three ways we can start making more ethical algorithms.)
  - Jeff Dean: AI isn't as smart as you think -- but it could be | TED Talk (18:22 What is AI, really? Jeff Dean, the head of Google's AI efforts, explains the underlying technology that enables artificial intelligence to do all sorts of things, from understanding language to diagnosing disease -- and presents a roadmap for building better, more responsible systems that have a deeper understanding of the world.)
- **Option 3:**
  - Yann LeCun: Deep learning, neural networks and the future of AI | TED Talk (55:02 Yann LeCun, the chief AI scientist at Facebook, helped develop the deep learning algorithms that power many artificial intelligence systems today. In conversation with head of TED Chris Anderson, LeCun discusses his current research into self-supervised

machine learning, how he's trying to build machines that learn with common sense (like humans) and his hopes for the next conceptual breakthrough in AI.)

Class 2: Image Recognition

- **Main Assignments: Do Them All**
  - How image recognition works: A complete guide to image recognition
  - Ferguson, Andrew Guthrie, "Facial Recognition and the Fourth Amendment" (2021). Minnesota Law Review. 3204. "Facial Recognition and the Fourth Amendment" by Andrew Guthrie Ferguson
    - Read pages 1109-1196.
- **Supplementary Assignment: Complete One Option**
  - **Option 1:**
    - Kade Crockford: What you need to know about face surveillance | TED Talk (12:40 Privacy isn't dead, but face surveillance technology might kill it, says civil rights advocate Kade Crockford. In an eye-opening talk, Kade outlines the startling reasons why this invasive technology -- powered by often-flawed facial recognition databases that track people without their knowledge -- poses unprecedented threats to your fundamental rights. Learn what can be done to ban government use before it's too late.)
    - Facial Recognition: What Happens When We're Tracked Everywhere We Go? - The New York Times
  - **Option 2:**
    - Joseph Redmon: How computers learn to recognize objects instantly | TED Talk (7: 29 Ten years ago, researchers thought that getting a computer to tell the difference between a cat and a dog would be almost impossible. Today, computer vision systems do it with greater than 99 percent accuracy. How? Joseph Redmon works on the YOLO (You Only Look Once) system, an open-source method of object detection that can identify objects in images and video -- from zebras to stop signs -- with lightning-quick speed.)
    - Google Photos Tags Two African-Americans As Gorillas Through Facial Recognition Software
    - When It Comes to Gorillas, Google Photos Remains Blind | WIRED
  - **Option 3:**
    - What facial recognition steals from us (9:42 Researchers at Facebook, Google, and other institutions have nearly perfected techniques for automated facial recognition. The result of that research is that your face isn't just a unique part of your body anymore, it's biometric data that can be copied an infinite number of times and stored forever. In this video, we explain how facial recognition technology works, where it came from, and what's at stake.)
    - Who thought it was a good idea to have facial recognition software?

Class 3: Text Recognition

- **Main Assignments: Do Them All**
  - A.I. Is Mastering Language. Should We Trust What It Says? - The New York Times
  - How NLP works: Natural Language Processing Explained Simply - HDS
  - Legal implications: NATURAL LANGUAGE PROCESSING
- **Supplementary Assignment: Complete One Option**
  - **Option 1:**
    - GPT-3, explained: OpenAI's new language AI is uncanny, funny- and a big deal - Vox
    - Meet GPT-3. It Has Learned to Code (and Blog and Argue). - The New York Times
  - **Option 2:**
    - Brookings: Detecting and mitigating bias in natural language processing
  - **Option 3:**
    - The (Un)ethical Story of GPT-3: OpenAI's Million Dollar Model | Matthew Burruss
    - AI game bans players for NSFW stories it generated itself • The Register

Class 4: Speech Recognition, Deepfakes, and Digital Assistants

- **Main Assignments: Do Them All**
  - Deepfakes and Ethics: https://aihalapathirana.medium.com/deepfakes-and-ethics-9b6656ddc07b
  - The Ethics of Smart Devices That Analyze How We Speak
  - Maurice E. Stucke & Ariel Ezrachi, *How Digital Assistants Can Harm Our Economy, Privacy, and Democracy*, 32 BERKELEY TECH. L.J. 1239 (2017). https://heinonline.org/HOL/P?h=hein.journals/berktech32&i=1283 (58 pages)
- **Supplementary Assignment: Complete One Option**
  - **Option 1:**
    - William Murphy, *Hey, Alexa.. I Take the Fifth*, 36 CRIM. Just. 24 (2021). https://heinonline.org/HOL/P?h=hein.journals/cjust36&i=99 (4 pages)
    - John G. Browning & Lisa Angelo, *Alexa, Testify*, 82 TEX. B.J. 506 (2019). https://heinonline.org/HOL/P?h=hein.barjournals/texbarj0082&i=506 (2 pages)
  - **Option 2:**
    - The Weekly | Deepfakes — Believe at Your Own Risk - The New York Times (26 minutes)
  - **Option 3:**
    - Police Want Your Smart Speaker—Here's Why | WIRED
  - **Option 4:**
    - Smart speakers offer new legal challenges as privacy goes public

        - Wired: The Race to Hide Your Voice

Class 5: Social Media and Entertainment

- **Main Assignments: Do Them All**
    - Watch The Social Dilemma | Netflix Official Site
    - Sinan Aral: How we can protect truth in the age of misinformation | TED Talk (14:55 Fake news can sway elections, tank economies and sow discord in everyday life. Data scientist Sinan Aral demystifies how and why it spreads so quickly -- citing one of the largest studies on misinformation -- and identifies five strategies to help us unweave the tangled web between true and false.)
    - An Ethics Perspective On Facebook
    - Solving the problem of racially discriminatory advertising on Facebook
- **Supplementary Assignment: Complete One Option**
    - **Option 1:**
        - Kevin Slavin: How algorithms shape our world | TED Talk (15:06 We live in a world run by algorithms, computer programs that make decisions or solve problems for us. In this riveting, funny talk, Kevin Slavin shows how modern algorithms determine stock prices, espionage tactics, even the movies you watch. But, he asks: If we depend on complex algorithms to manage our daily decisions -- when do we start to lose control?)
        - Rachel Casey, *John Stuart Mill and Social Media: Evaluating the Ethics of De-Platforming*, 4 U. CENT. FLA. DEP't LEGAL Stud. L.J. 15 (2021). https://heinonline.org/HOL/P?h=hein.journals/ucflaegs4&i=16 (19 pages)
    - **Option 2**:
        - Kathleen McGarvey Hidy, *The Speech Gods: Freedom of Speech, Censorship, and Cancel Culture in the Age of Social Media*, 61 Washburn L.J. 99 (2021). https://heinonline.org/HOL/P?h=hein.journals/wasbur61&i=107 (56 pages)
    - **Option 3:**
        - Kathleen McGarvey Hidy, *Social Media Use and Viewpoint Discrimination: A First Amendment Judicial Tightrope Walk with Rights and Risks Hanging in the Balance*, 102 MARQ. L. REV. 1045 (2019). https://heinonline.org/HOL/P?h=hein.journals/marqlr102&i=1079 (40 pages)

Class 6: Credit Scores

- **Main Assignments: Do Them All**
    - Janine S. Hiller, *Fairness in the Eyes of the Beholder: AI; Fairness; and Alternative Credit Scoring*, 123 W. VA. L. REV. 907 (2021). https://heinonline.org/HOL/P?h=hein.journals/wvb123&i=937 (27 pages)

    - Yu-Jie Chen, Ching-Fu Lin & Han-Wei Liu, Rule of Trust: The Power and Perils of China's Social Credit Megaproject, 32 COLUM. J. Asian L. 1 (2018). https://heinonline.org/HOL/P?h=hein.journals/colas32&i=3 (32 pages)
- **Supplementary Assignment: Complete One Option**
    - **Option 1:**
        - Brookings: Reducing bias in AI-based financial services
    - **Option 2:**
        - Stanford: How Flawed Data Aggravates Inequality in Credit
    - **Option 3:**
        - HBR: AI Can Make Bank Loans More Fair
    - **Option 4:**
        - Brookings: Credit denial in the age of AI

Class 7: Job Hiring and the Workforce

- **Main Assignments: Do Them All**
    - Pauline T. Kim & Matthew T. Bodie, *Artificial Intelligence and the Challenges of Workplace Discrimination and Privacy*, 35 A.B.A. J. LAB. & EMP. L. 289 (2021). https://heinonline.org/HOL/P?h=hein.journals/lablaw35&i=305 (27 pages)
    - How AI is Deciding Who Gets Hired (15:23)
    - Bloomberg: How AI Is Deciding Who Gets Hired
    - HBR: The Legal and Ethical Implications of Using AI in Hiring
- **Supplementary Assignment: Complete One Option**
    - **Option 1:**
        - How Job Applicants Try to Hack Résumé-Reading Software | WIRED
        - Kevin Roose: The value of your humanity in an automated future | TED Talk (10:40 To futureproof your job against robots and AI, you should learn how to code, brush up on your math skills and crack open an engineering textbook, right? Wrong. In this surprisingly comforting talk, tech journalist Kevin Roose makes the case that rather than trying to compete with the machines, we should instead focus on what makes us uniquely human.)
    - **Option 2:**
        - Will robots spreading through workforce crowd out humans? | AP News
        - Anthony Goldbloom: The jobs we'll lose to machines -- and the ones we won't | TED Talk (4:27 Machine learning isn't just for simple tasks like assessing credit risk and sorting mail anymore -- today, it's capable of far more complex applications, like grading essays and diagnosing diseases. With these advances comes an uneasy question: Will a robot do your job in the future?)
    - **Option 3:**
        - Part 1: How many jobs do robots really replace? | MIT News

        - Part 2: Robots help some firms, even while workers across industries struggle | MIT News
        - Part 3: Study finds stronger links between automation and inequality | MIT News

Class 8: Finance

- **Main Assignments: Do Them All**
    - Hilary J. Allen, *Driverless Finance*, 10 HARV. Bus. L. REV. 157 (2020). https://heinonline.org/HOL/P?h=hein.journals/hbusrew10&i=163 (47 pages)
    - William Magnuson, *Artificial Financial Intelligence*, 10 HARV. Bus. L. REV. 337 (2020). https://heinonline.org/HOL/P?h=hein.journals/hbusrew10&i=345 (43 pages)
- **Supplementary Assignment: Complete One Option**
    - **Option 1:**
        - Tom C. W. Lin, *Artificial Intelligence, Finance, and the Law*, 88 FORDHAM L. REV. 531 (2019). https://heinonline.org/HOL/P?h=hein.journals/flr88&i=547 (21 pages)
    - **Option 2:**
        - Kristin Johnson, Frank Pasquale & Jennifer Chapman, Artificial Intelligence, Machine Learning, and Bias in Finance: Toward Responsible Innovation, 88 FORDHAM L. REV. 499 (2019). https://heinonline.org/HOL/P?h=hein.journals/flr88&i=515 (31 pages)
    - **Option 3:**
        - The Future Robo Adviser: Smart and Ethical? - WSJ
        - Brookings: Can AI model economic choices?
        - Brookings: Reducing bias in AI-based financial services

Class 9: War

- **Main Assignments: Do Them All**
    - Daniel Suarez: The kill decision shouldn't belong to a robot | TED Talk (13:07 As a novelist, Daniel Suarez spins dystopian tales of the future. But on the TEDGlobal stage, he talks us through a real-life scenario we all need to know more about: the rise of autonomous robotic weapons of war. Advanced drones, automated weapons and AI-powered intelligence-gathering tools, he suggests, could take the decision to make war out of the hands of humans.)
    - Brookings: Democratizing harm: Artificial intelligence in the hands of nonstate actors (19 pages)
    - Brookings: Artificial intelligence in war: Human judgment as an organizational strength and a strategic liability (12 pages)
- **Supplementary Assignment: Complete One Option**
    - **Option 1:**

        - Joaquin Rodriguez Alvarez, *Social Challenges of Artificial Intelligence: The Case of Lethal Autonomous Systems*, 2018 J. LAW 244 (2018). https://heinonline.org/HOL/P?h=hein.journals/tsujrnl2018&i=244 (25 pages)
    - **Option 2:**
        - Brian Seamus Haney, *Applied Artificial Intelligence in Modern Warfare and National Security Policy*, 11 Hastings Sci. & TECH. L.J. 61 (2020). https://heinonline.org/HOL/P?h=hein.journals/hascietlj11&i=62 (25 pages – Don't read the Deep Reinforcement section unless you want to. You won't be quizzed on the math.)
    - **Option 3:**
        - Ashley Deeks, Noam Lubell & Daragh Murray, *Machine Learning, Artificial Intelligence, and the Use of Force by States*, 10 J. NAT'l Sec. L. & POL'y 1 (2019). https://heinonline.org/HOL/P?h=hein.journals/jnatselp10&i=6 (25 pages)

Class 10: Autonomous Vehicles

- **Main Assignments: Do Them All**
    - How driverless cars will change our world - BBC Future
    - https://podcasts.google.com/feed/aHR0cDovL2lxMi5pcTJ1cy5saWJzeW5wcm8uY29tL3Jzcw/episode/OGIxOTAwMTAtMWQ5OS0xMWU4LWI1Y2YtMmI2MzE5ZjU2MzU3?ep=14
    - Michael Mattioli, *Autonomy in the Age of Autonomous Vehicles*, 24 B.U. J. Sci. & TECH. L. 277 (2018). https://heinonline.org/HOL/P?h=hein.journals/jstl24&i=299 (21 pages)
    - MIT Moral Machine questionnaire: https://www.moralmachine.net/
    - The Moral Machine experiment
    - Driverless autonomous trucks lead the way | Deloitte Insights
- **Supplementary Assignment: Complete One Option**
    - **Option 1:**
        - What moral decisions should driverless cars make? | Iyad Rahwan (13:36)
    - **Option 2:**
        - The ethical dilemma of self-driving cars - Patrick Lin (4:16)
        - What will self-driving trucks mean for truck drivers? - BBC News
    - **Option 3:**
        - Would you sacrifice one person to save five? - Eleanor Nelsen (4:56)
        - How do self-driving cars "see"? - Sajan Saini (5:25)

Class 11: Medical Care

- **Main Assignments: Do Them All**
  - These Algorithms Look at X-Rays—and Somehow Detect Your Race | WIRED
  - Here's how an algorithm guides a medical decision
  - Jim Collins: How we're using AI to discover new antibiotics | TED Talk (7:07)
  - Sharona Hoffman & Andy Podgurski, *Artificial Intelligence and Discrimination in Health Care*, 19 YALE J. HEALTH POL'y L. & Ethics 1 (2020). https://heinonline.org/HOL/P?h=hein.journals/yjhple19&i=490 (44 pages)
- **Supplementary Assignment: Complete One Option**
  - **Option 1:**
    - Barry R. Furrow, *The Limits of Current A.I. in Health Care: Patient Safety Policing in Hospitals*, 12 N.E. U. L.R. 1 (2020). https://heinonline.org/HOL/P?h=hein.journals/norester12&i=16 (54 pages)
  - **Option 2:**
    - Sahar Takshi, *Unexpected Inequality: Disparate-Impact from Artificial Intelligence in Healthcare Decisions*, 34 J.L. & HEALTH 215 (2021). https://heinonline.org/HOL/P?h=hein.journals/jlah34&i=222 (32 pages)
    - Risks and remedies for artificial intelligence in health care
  - **Option 3:**
    - Watson: The Promise: Paging Dr. Watson: Artificial Intelligence As a Prescription for Health Care | WIRED
    - Watson: The Failure: What Ever Happened to IBM's Watson? - The New York Times
    - AI helping with breast cancer:
      - MIT professor's AI predicts breast cancer risk from mammograms - The Washington Post
      - Artificial Intelligence Makes Bad Medicine Even Worse | WIRED
    - AlphaFold: AlphaFold Is The Most Important Achievement In AI—Ever

Class 12: AI as Judge

- **Main Assignments: Do Them All**
  - Sandra G. Mayson, *Bias in, Bias out*, 128 YALE L. J. 2218 (2019). Bias In, Bias Out (75 pages)
  - Understanding risk assessments: Understanding risk assessment instruments in criminal justice
- **Supplementary Assignment: Complete One Option**
  - **Option 1:**
    - AI poses risks in criminal justice: Artificial intelligence poses serious risks in the criminal justice system - The Johns Hopkins News-Letter.
    - Bob Lambrechts, *May It Please the Algorithm: The Future of A.I. in the Practice of Law*, 89 J. KAN. B. Ass'n 36 (2020).

https://heinonline.org/HOL/P?h=hein.barjournals/jkabr0089&i=41 (6 pages)

- **Option 2:**
  - Tania Sourdin, *Judge v Robot? Artificial Intelligence and Judicial Decision-Making*, 41 U.N.S.W.L.J. 1114 (2018). https://heinonline.org/HOL/P?h=hein.journals/swales41&i=1130 (19 pages)

Class 13: AI as Board Member

- **Main Assignments: Do Them All**
  - Board of Directors Explained:
    - What is a Board of Directors? - The Corporate Governance Institute
    - What is a board of directors and what does it do?
  - Akshaya Kamalnath, *The Perennial Quest for Board Independence: Artificial Intelligence to the Rescue?*, 83 ALB. L. REV. 43 (2019). https://heinonline.org/HOL/P?h=hein.journals/albany83&i=51 (18 pages)
  - Gramitto Ricci, Sergio Alberto, Artificial Agents in Corporate Boardrooms (August 20, 2020). Cornell Law Review, Vol. 105, No. 3, 2020, Available at SSRN: Artificial Agents in Corporate Boardrooms (40 pages)
- **Supplementary Assignment: Complete One Option**
  - **Option 1:**
    - Why Not Appoint an Algorithm to Your Corporate Board?
  - **Option 2:**
    - AI in the Boardroom: The Next Realm of Corporate Governance
  - **Option 3:**
    - AI in the Boardroom: A Revolutionary Construct? – The FinReg Blog

Class 14: Artificial General Intelligence

- **Main Assignments: Do Them All**
  - David Atkinson, *AI for President?*, 3 THE JOURNAL of ROBOTICS, ARTIFICIAL INTELLIGENCE & LAW (Fastcase) 89 (2020). https://heinonline.org/HOL/P?h=hein.journals/rail3&i=91 (30 pages)
  - OpenAI charter: OpenAI Charter
  - Sam Harris: Can we build AI without losing control over it? | TED Talk (14:19)
- **Supplementary Assignment: Complete One Option**
  - **Option 1:**
    - Janelle Shane: The danger of AI is weirder than you think | TED Talk (10:20)
    - Nick Bostrom: What happens when our computers get smarter than we are? | TED Talk (16:22)
    - Grady Booch: Don't fear superintelligent AI | TED Talk (10:12)

- **Option 2:**
  - The Artificial Intelligence Revolution: Part 1 - Wait But Why
  - https://waitbutwhy.com/2015/01/artificial-Intelligence-Revolution-2.html